

Hello Soft Clustering (GMM)

T1. Using 3 mixtures, initialize your Gaussian with means (3,3), (2,2), and (-3,-3), and standard Covariance, I, the identity matrix. Use equal mixture weights as the initial weights. Repeat three iterations of EM. Write down $w_{n,j}$, μ_j , Σ_j (Sigma_) for each EM iteration. (You may do the calculations by hand or write code to do so)

$$w_{n,j} = \frac{p(x_n; \hat{\mu}_j^*, \hat{\Sigma}_j) m_j}{\sum_j p(x_n; \hat{\mu}_j^*, \hat{\Sigma}_j) m_j} \quad (1)$$

$w_{n,j}$ means the probability that data point n comes from Gaussian number j .

Maximization: Update the model parameters, ϕ , $\hat{\mu}_j^*$, $\hat{\Sigma}_j$.

$$m_j = \frac{1}{N} \sum_n w_{n,j} \quad (2)$$

$$\hat{\mu}_j^* = \frac{\sum_n w_{n,j} x_n}{\sum_n w_{n,j}} \quad (3)$$

$$\hat{\Sigma}_j = \frac{\sum_n w_{n,j} (x_n - \hat{\mu}_j^*)(x_n - \hat{\mu}_j^*)^T}{\sum_n w_{n,j}} \quad (4)$$

The above equation is used for full covariance matrices. For our small toy example, we will use diagonal covariance matrices, which can be acquired by setting the off-diagonal values to zero. In other words, $\Sigma_{(i,j)} = 0$, for $i \neq j$.

TODO: Complete functions below including

- Fill relevant parameters in each function.
- Implement computation and return values.

These functions will be used in T1-4.

```
In [1]:
import numpy as np
import matplotlib.pyplot as plt

# Hint: You can use this function to get gaussian distribution.
# https://docs.scipy.org/doc/scipy/reference/generated/scipy.stats.multivariate_normal.html
from scipy.stats import multivariate_normal

In [12]:
multivariate_normal.pdf([[1, 2], [3, mean=[1, 1], cov=[[1, 0], [0, 1]]])

Out[12]:
array([0.09653235, 0.09653235, 0.09653235])

In [13]:
class GMM:
    def __init__(self, mixture_weight, mean_params, cov_params):
        """
        Initialize GMM.
        """
        # Copy construction values.
        self.mixture_weight = mixture_weight
        self.mean_params = mean_params
        self.cov_params = cov_params

        # Initialize iteration.
        self.n_iter = 0

    def estimation_step(self, data):
        """
        TODO: Perform estimation step. Then, return w_{n,j} in eq. 1
        """

        # INSERT CODE HERE

        w = np.array([multivariate_normal.pdf(data, mean=self.mean_params[i], cov=self.cov_params[i]) * self.
            mixture_weight for i in range(len(self.mixture_weight))])

        print(w.shape)
        print(w)
        print(w.sum(axis=1))
        print(np.array([w.sum(axis=1) * 3]).T)

        w = w / w.sum(axis=1)[:, None] # norm probs across mixtures

        return w # assigned soft labels (data points, mixtures)

    def maximization_step(self, data, w):
        """
        TODO: Perform maximization step.
        (Update parameters in this GMM model.)
        """

        # INSERT CODE HERE

        self.mixture_weight = w.sum(axis=0) / len(data)

        self.mean_params = np.matmul(w.T, data) / w.sum(axis=0)[:, None]

        self.cov_params = np.array([np.matmul((data - self.mean_params[i]).T, (data - self.mean_params[i]) *
            self.cov_params[i] * np.eye(self.cov_params[i].shape[0])) for i in range(len(self.mixture_weight))])

    def get_log_likelihood(self, data):
        """
        TODO: Compute Log Likelihood.
        """

        # INSERT CODE HERE

        log_prob = np.log(np.array([multivariate_normal.pdf(data, mean=self.mean_params[i], cov=self.cov_params[i])
            for i in range(len(self.mixture_weight))]))

        return log_prob

    def print_iteration(self):
        """
        Print iteration information.
        """
        print("mixture_weight")
        print("mu")
        print("covariance matrix")
        print("log_prob")

    def perform_estimation_and_maximization_steps(self, data, num_iteations):
        """
        Perform estimation and maximization steps with num_iteations.
        Then, return list of log_likelihood from those iterations.
        """
        log_prob_list = []

        # Display initialization.
        if display:
            print("Initialization")
            self.print_iteration()

        for n_iter in range(num_iteations):
            # TODO: Perform EM step.

            # INSERT CODE HERE

            w = self.estimation_step(data)
            self.maximization_step(data, w)

            self.print_iteration()

            # Calculate log prob.
            log_prob = self.get_log_likelihood(data)
            log_prob_list.append(log_prob)

            # Display each iteration.
            if display:
                print(f"Iteration: {n_iter}")
                self.print_iteration()

        return log_prob_list

In [41]:
num_iteations = 3
num_mixture = 3
mixture_weight = [1] * num_mixture # m
mean_params = np.array([3,3], [2,2], [-3,-3]), dtype = float)
cov_params = np.array([np.eye(2)] * num_mixture)

X, y = np.array([3, 3, 2, 8, 6, 7, -3, -2, -7]), np.array([2, 3, 2, 8, 6, 7, -3, -4, -7])
data = np.vstack([X, y]).T

gmm = GMM(mixture_weight, mean_params, cov_params)

log_prob_list = gmm.perform_em_iteations(data, num_iteations)

Initialization
m :
[[ 3.  3.]
 [ 2.  2.]
 [-3. -3.]]
covariance matrix :
[[[1. 0.]
 [0. 1.]]
 [[1. 0.]
 [0. 1.]]
 [[1. 0.]
 [0. 1.]]]

-----
m :
[0.45757242 0.20909425 0.33333333]
mu :
[[ 5.70992692  5.81887265]
 [ 1.67718211  2.14523106]
 [-4.        -4.66666666]]
covariance matrix :
[[[4.53619412 0.        ]
 [0.        4.28780611]]
 [[0.51645579 0.        ]
 [0.        0.13152618]]
 [[4.66666668 0.        ]
 [0.        2.88888891]]]

Iteration: 0
m :
[0.45757242 0.20909425 0.33333333]
mu :
[[ 5.70992692  5.81887265]
 [ 1.67718211  2.14523106]
 [-4.        -4.66666666]]
covariance matrix :
[[[4.53619412 0.        ]
 [0.        4.28780611]]
 [[0.51645579 0.        ]
 [0.        0.13152618]]
 [[4.66666668 0.        ]
 [0.        2.88888891]]]

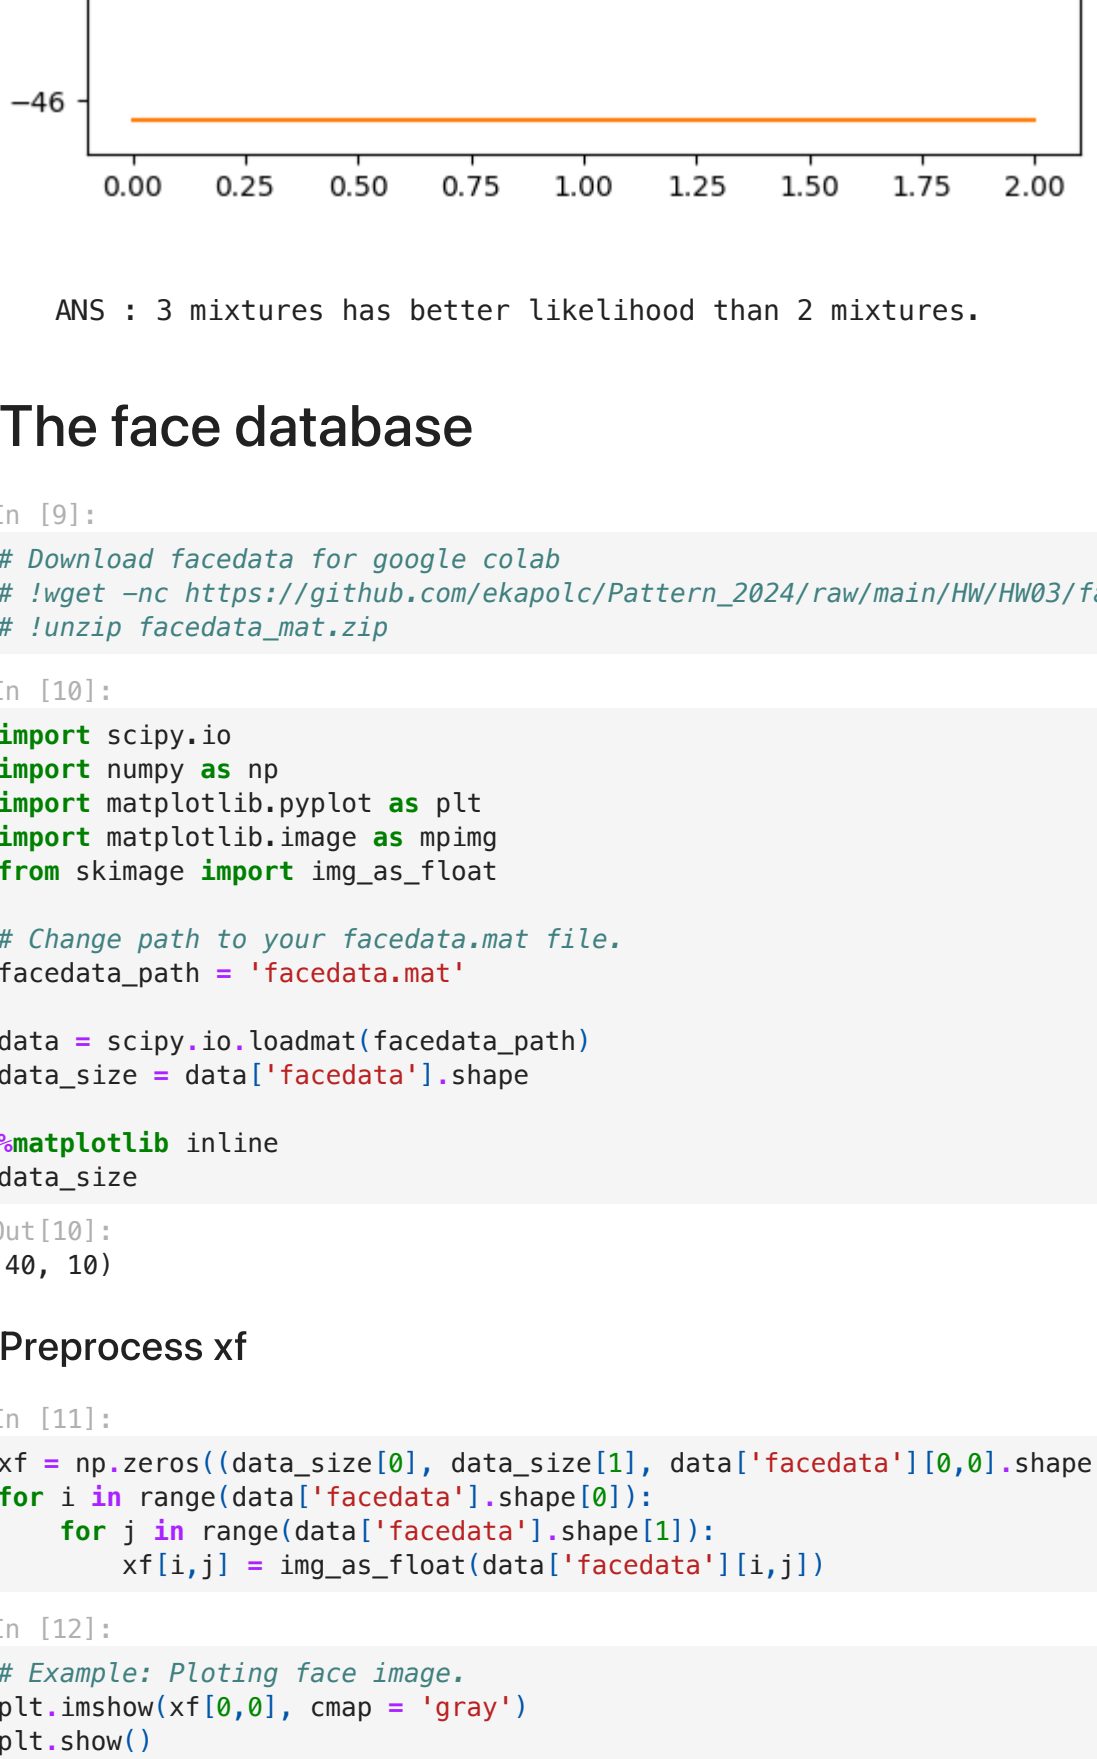
-----
m :
[0.40711618 0.25954961 0.33333421]
mu :
[[ 6.27176215  6.27262711]
 [ 1.72091544  2.14764012]
 [-3.99998589 -4.6666488 ]]
covariance matrix :
[[[2.94672736 0.        ]
 [0.        2.93847196]]
 [[0.49649261 0.        ]
 [0.        0.12584815]]
 [[4.66673088 0.        ]
 [0.        2.88900236]]]

-----
Iteration: 1
m :
[0.40711618 0.25954961 0.33333421]
mu :
[[ 6.27176215  6.27262711]
 [ 1.72091544  2.14764012]
 [-3.99998589 -4.6666488 ]]
covariance matrix :
[[[2.94672736 0.        ]
 [0.        2.93847196]]
 [[0.49649261 0.        ]
 [0.        0.12584815]]
 [[4.66673088 0.        ]
 [0.        2.88900236]]]

-----
m :
[0.36070909 0.30595677 0.33333414]
mu :
[[ 6.6962644  6.69629468]
 [ 1.91071238  2.27383436]
 [-3.99998073 -4.666591 ]]
covariance matrix :
[[[1.73961067 0.        ]
 [0.        1.73929602]]
 [[0.62898406 0.        ]
 [0.        0.1988491 ]]
 [[4.66672942 0.        ]
 [0.        2.88899545]]]

-----
Iteration: 2
m :
[0.36070909 0.30595677 0.33333414]
mu :
[[ 6.6962644  6.69629468]
 [ 1.91071238  2.27383436]
 [-3.99998073 -4.666591 ]]
covariance matrix :
[[[1.73961067 0.        ]
 [0.        1.73929602]]
 [[0.62898406 0.        ]
 [0.        0.1988491 ]]
 [[4.66672942 0.        ]
 [0.        2.88899545]]]
```

T2. Plot the log likelihood of the model given the data after each EM step. In other words, plot $\log \prod_n p(\text{vec}(x_n) | \phi, \text{vec}(\mu_j, \Sigma_j))$. Does it goes up every iteration just as we learned in class?



T3. Using 2 mixtures, initialize your Gaussian with means (3,3) and (-3,-3), and standard Covariance, I, the identity matrix. Use equal mixture weights as the initial weights. Repeat three iterations of EM. Write down $w_{n,j}$, μ_j , Σ_j (Sigma_) for each EM iteration.

```
In [6]:
num_mixture = 2
mixture_weight = [1] * num_mixture
mean_params = np.array([3,3], [-3,-3]), dtype = float)
cov_params = np.array([np.eye(2)] * num_mixture)

# INSERT CODE HERE
gmm2 = GMM(mixture_weight, mean_params, cov_params)
log_prob_list2 = gmm2.perform_em_iteations(data, num_iteations)

Initialization
m :
[[ 3.  3.]
 [-3. -3.]]
covariance matrix :
[[[1. 0.]
 [0. 1.]]
 [[1. 0.]
 [0. 1.]]]

-----
m :
[0.66666666 0.33333334]
mu :
[[ 4.50000001  4.66666667]
 [-3.99999997 -4.66666663]]
covariance matrix :
[[[6.91944755 0.        ]
 [0.        5.88888889]]
 [[4.66666677 0.        ]
 [0.        2.88888891 ]]]

Iteration: 0
m :
[0.66666666 0.33333334]
mu :
[[ 4.50000001  4.66666667]
 [-3.99999997 -4.66666663]]
covariance matrix :
[[[6.91944755 0.        ]
 [0.        5.88888889]]
 [[4.66666677 0.        ]
 [0.        2.88888891 ]]]

-----
m :
[0.66669436 0.33330564]
mu :
[[ 4.49961311  4.66620178]
 [-3.99993241 -4.66652311]]
covariance matrix :
[[[6.91944755 0.        ]
 [0.        5.89275124]]
 [[4.66806942 0.        ]
 [0.        2.89103318]]]

Iteration: 1
m :
[0.66669436 0.33330564]
mu :
[[ 4.49961311  4.66620178]
 [-3.99993241 -4.66652311]]
covariance matrix :
[[[6.91944755 0.        ]
 [0.        5.89275124]]
 [[4.66806942 0.        ]
 [0.        2.89103318]]]

-----
m :
[0.66669453 0.33330547]
mu :
[[ 4.49961884  4.66619083]
 [-3.99993206 -4.66651141]]
covariance matrix :
[[[6.91946372 0.        ]
 [0.        5.8927741 ]]
 [[4.66807754 0.        ]
 [0.        2.89104566]]]

Iteration: 2
m :
[0.66669453 0.33330547]
mu :
[[ 4.49961884  4.66619083]
 [-3.99993206 -4.66651141]]
covariance matrix :
[[[6.91946372 0.        ]
 [0.        5.8927741 ]]
 [[4.66807754 0.        ]
 [0.        2.89104566]]]
```

T4. Plot the log likelihood of the model given the data after each EM step. Compare the log likelihood between using two mixtures and three mixtures. Which one has the better likelihood?



The face database

```
In [91]:
# Download facedata for google colab
# wget -nc https://github.com/ekapolc/Pattern_2024/raw/main/HM/HM03/facedata_mat.zip
# unzip facedata_mat.zip
```

```
In [10]:
import scipy.io
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
from skimage import io as imgio
```

```
# Change path to your facedata.mat file.
facedata_path = 'facedata.mat'

data = scipy.io.loadmat(facedata_path)
data_size = data['facedata'].shape
```

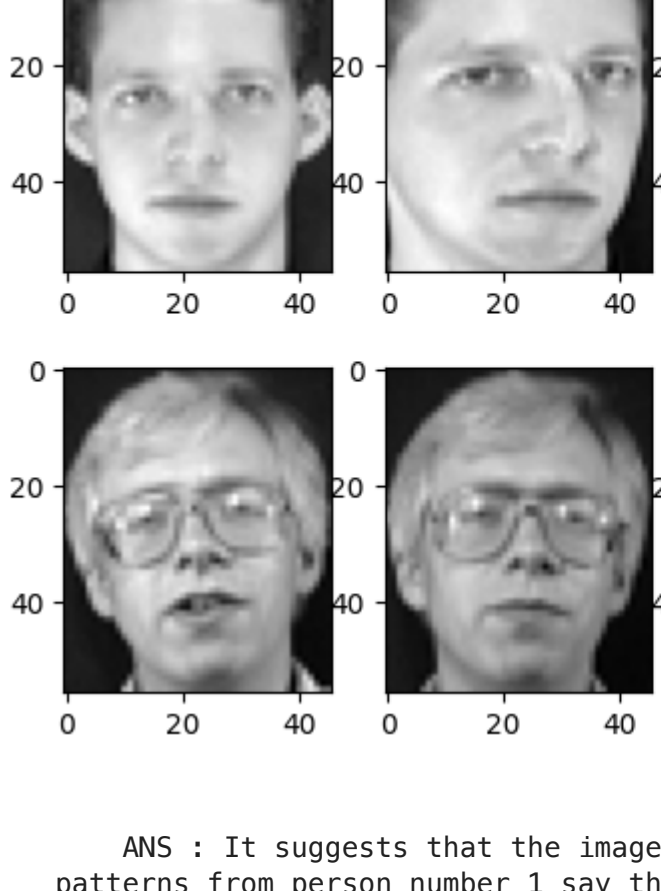
```
matplotlib inline
data_size
```

```
Out[10]:
(40, 10)
```

Preprocess xf

```
xf = np.zeros([data_size[0], data_size[1], data['facedata'][0,0].shape[0], data['facedata'][0,0].shape[1]])
for i in range(data['facedata'].shape[0]):
    for j in range(data['facedata'].shape[1]):
        xf[i,j] = imgio_float(data['facedata'][i,j])
```

```
In [12]:
# Example: Plotting face image.
plt.imshow(xf[0,0], cmap='gray')
plt.show()
```



T5. What is the Euclidean distance between $xf[0,0]$ and $xf[0,1]$? What is the Euclidean distance between $xf[0,0]$ and $xf[1,0]$? Does the numbers make sense? Do you think these numbers will be useful for face verification?

```
In [13]:
def L2_dist(x1, x2):
    """
    TODO: Calculate L2 distance.
    """
    return np.sqrt((x1 - x2) ** 2).sum())

# Test L2_dist
def test_L2_dist():
    test_A = np.array([1, 2, 3, 4])
    test_B = np.array([1, 2, 1, 5, 6], [7, 8])
    expected_matrix = np.sqrt(np.array([[0, 32, 72], [8, 8, 32]]))
    assert (generate_similarity_matrix(test_A, test_B) == expected_matrix).all()
```

```
test_generate_similarity_matrix()
```

```
In [16]:
# TODO: Show similarity matrix between T and D.
```

```
# INSERT CODE HERE
T = xf[:, :3]
D = xf[:, 3:]

similarity_matrix = generate_similarity_matrix(organize_shape(T), organize_shape(D))
```

```
In [17]:
plt.imshow(similarity_matrix, cmap='gray')
```

```
Out[17]:
<matplotlib.image.AxesImage at 0x169a957f10>
```


ANS : I think the numbers make sense because $xf[0,0]$ and $xf[1,0]$ is the same pose that have the similar ratio of face and background, but $xf[0,0]$ and $xf[0,1]$ the pose is different so there're more distance between them. But I think it's not useful for face verification because the distance within the same person should be smaller than the distance between different person to tell that the image shown is the same person or not.

T6. Write a function that takes in a set of feature vectors T and a set of feature vectors D, and then output the similarity matrix A. Show the matrix as an image. Use the feature vectors from the first 3 images from all 40 people for list T (in order $x[0,0]$, $x[0,1]$, $x[0,2]$, $x[1,0]$, $x[1,1]$, ..., $x[39,2]$). Use the feature vectors from the remaining 7 images from all 40 people for list D (in order $x[0,3]$, $x[0,4]$, $x[0,5]$, $x[1,6]$, $x[0,6]$, $x[0,7]$, $x[0,8]$, $x[0,9]$, $x[1,9]$, $x[1,4]$, ..., $x[39,9]$). We will treat T as our training images and D as our testing images

```
In [18]:
def organize_shape(matrix):
    """
    TODO (Optional): Reduce matrix dimension of 2D image to 1D and merge people and image dimension.
    This function can be useful at organizing matrix shapes.
    """
    Example:
    Input shape: (people_index, image_index, image_shape[0], image_shape[1])
    Output shape: (people_index*image_index, image_shape[0]*image_shape[1])
    """
    return matrix.reshape(matrix.shape[0] * matrix.shape[1], matrix.shape[2] * matrix.shape[3])
```

```
def generate_similarity_matrix(A, B):
    """
    TODO: Calculate similarity matrix M.
    Which  $M[i, j]$  is a distance between  $A[i]$  and  $B[j]$ .
    """
    # INSERT CODE HERE
    similarity_matrix = np.array([[L2_dist(A[i], B[j]) for j in range(len(B)) for i in range(len(A))])
    return similarity_matrix
```

```
def test_generate_similarity_matrix():
    test_A = np.array([1, 2, 3, 4])
    test_B = np.array([1, 2, 1, 5, 6], [7, 8])
    expected_matrix = np.sqrt(np.array([[0, 32, 72], [8, 8, 32]]))
    assert (generate_similarity_matrix(test_A, test_B) == expected_matrix).all()
```

```
test_generate_similarity_matrix()
```

```
In [19]:
# TODO: Show similarity matrix between T and D.
```

```
# INSERT CODE HERE
fig, axes = plt.subplots(2, 5, figsize=(10, 5))
for i in range(5):
    axes[0, i].imshow(xf[0, i], cmap='gray')
    axes[1, i].imshow(xf[1, i], cmap='gray')
```

```
Out[19]:
<matplotlib.image.AxesImage at 0x169a957f10>
```

ANS : It suggests that the images from person number 2 are similar to each other. The patterns from person number 1 say that the images from person 1 are different from each other.

T8. Write a function that takes in the similarity matrix created from the previous part, and a threshold as inputs. The outputs of the function are the true positive rate and the false alarm rate of the face verification task (280 Test images, tested on 40 people, a total of 1000 testing per threshold). What is the true positive rate and the false alarm rate for $t = 10$?

```
In [64]:
similarity_matrix.shape
```

```
Out[64]:
(120, 280)
```

```
In [65]:
120/3=280
```

```
Out[65]:
280
```

```
In [67]:
280 = 40
```

```
Out[67]:
1200
```



```
In [79]:
def evaluate_performance(similarity_matrix, threshold):
    """
    T000: Calculate true positive rate and false alarm rate from given similarity_matrix and threshold.
    """
    # INSERT CODE HERE
    """
    labels = np.array([np.arange(1, 41)] * (7 * 40)).flatten()
    labels = np.arange(1, 41).repeat(7*40)
    print(labels)
    """

    labels = []
    predictions = []
    for i in range(0, similarity_matrix.shape[0], 3):
        for j in range(similarity_matrix.shape[1]):
            # labels.append(i // 3 + 1)
            # predictions.append(i // 3 + 1)
            labels.append(j // 7 + 1)
            # print(similarity_matrix[i : i + 3, j])
            predictions.append(np.where(similarity_matrix[i:i+3, j].min() < threshold, 1 // 3+1, 0))

    labels = np.array(labels)
    # print(labels)
    # print(predictions)
    predictions = np.array(predictions)

    # import pandas as pd
    # print(pd.DataFrame(predictions).value_counts())
    # print(predictions.shape)
    # print(predictions)
    true_pos = np.sum((predictions == labels) & (predictions != 0))
    false_pos = np.sum((predictions == labels) & (predictions != 0))
    true_neg = np.sum((predictions == 0) & (predictions != labels))
    false_neg = np.sum((predictions == 0) & (labels != 0))
    # print(true_pos, false_pos, true_neg, false_neg)
    true_pos_rate = true_pos / (true_pos + false_neg)
    false_pos_rate = false_pos / (false_pos + true_neg)
    return true_pos_rate, false_pos_rate

# Quick check
# (true_pos_rate, false_neg_rate) should be (0.9928571428571429, 0.33587326007326005)
evaluate_performance(similarity_matrix, 9.5)

Out[79]:
(0.036865137249701634, 0.3350119625819447)

In [80]:
# INSERT CODE HERE
evaluate_performance(similarity_matrix, 10)

Out[80]:
(0.044884169884169885, 0.4563684644263346)

ANS: (0.044884169884169885, 0.4563684644263346)
```

T9. Plot the ROC curve for this simple verification system. What should be the minimum threshold to generate the ROC curve? What should be the maximum threshold? Your RoC should be generated from at least 1000 threshold levels equally spaced between the minimum and the maximum. (You should write a function for this).

```
In [81]:
def calculate_roc(input_mat):
    """
    T000: Calculate a list of true_pos_rate and a list of false_neg_rate from the given matrix.
    """
    # INSERT CODE HERE
    thresholds = np.linspace(input_mat.min(), input_mat.max(), 1000)

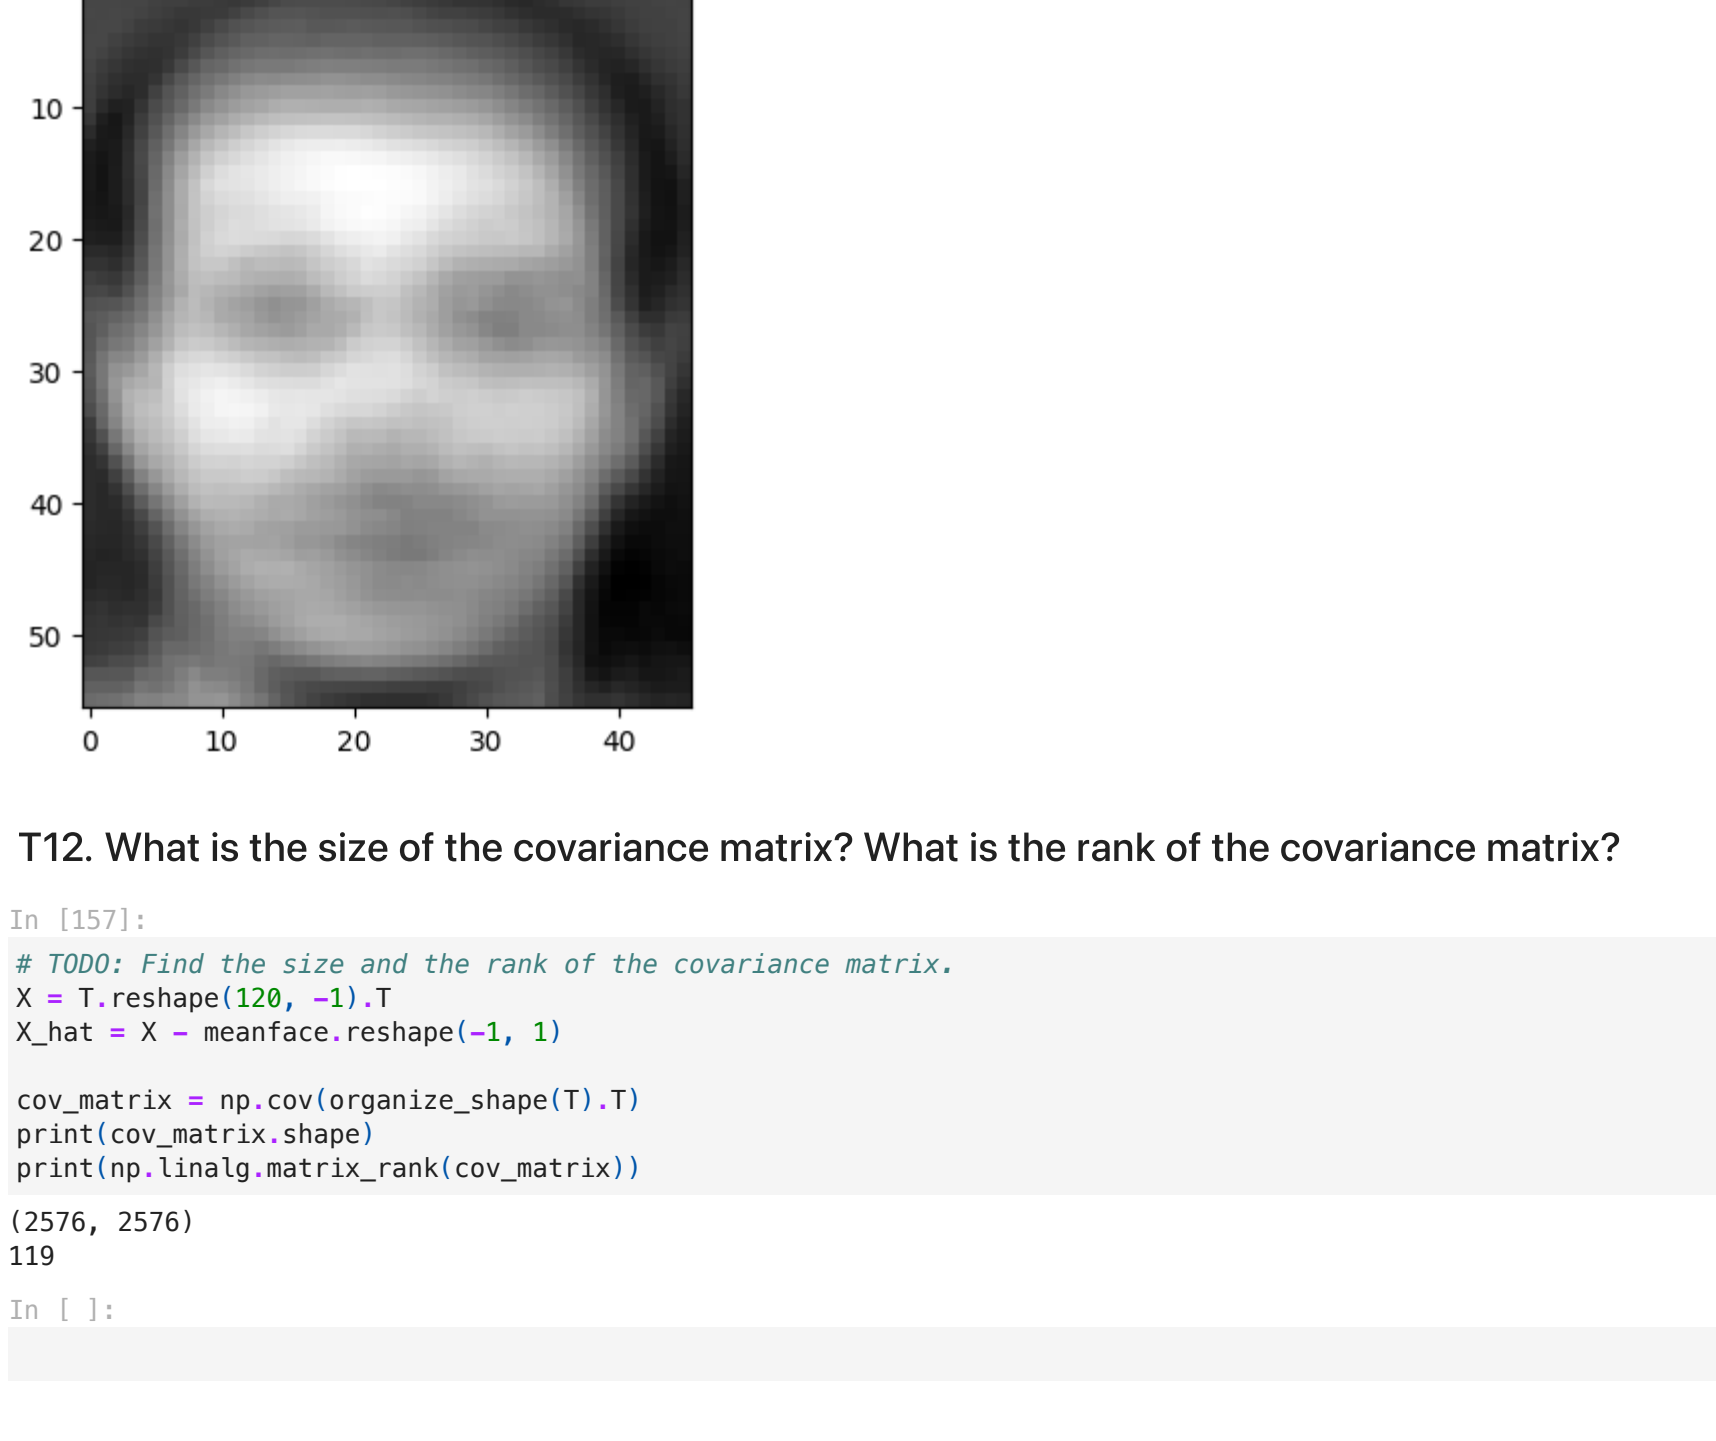
    tpr_list, far_list = [], []

    for threshold in thresholds:
        tpr, far = evaluate_performance(similarity_matrix, threshold)
        tpr_list.append(tpr)
        far_list.append(far)

    return tpr_list, far_list

def plot_roc(input_mat):
    """
    T000: Plot RoC Curve from a given matrix.
    """
    # INSERT CODE HERE
    plt.plot(calculate_roc(input_mat)[1], calculate_roc(input_mat)[0])
    plt.xlabel("False Alarm Rate")
    plt.ylabel("True Positive Rate")

In [82]:
# INSERT CODE HERE
plot_roc(similarity_matrix)
```



ANS: The minimum threshold should be smallest value of the similarity matrix. The maximum threshold should be the largest value of the similarity matrix.

T10. What is the EER (Equal Error Rate)? What is the recall rate at 0.1% false alarm rate? (Write this in the same function as the previous question)

```
In [83]:
def evaluate_eer(similarity_matrix, threshold):
    pass

evaluate_eer(similarity_matrix, 9.5)

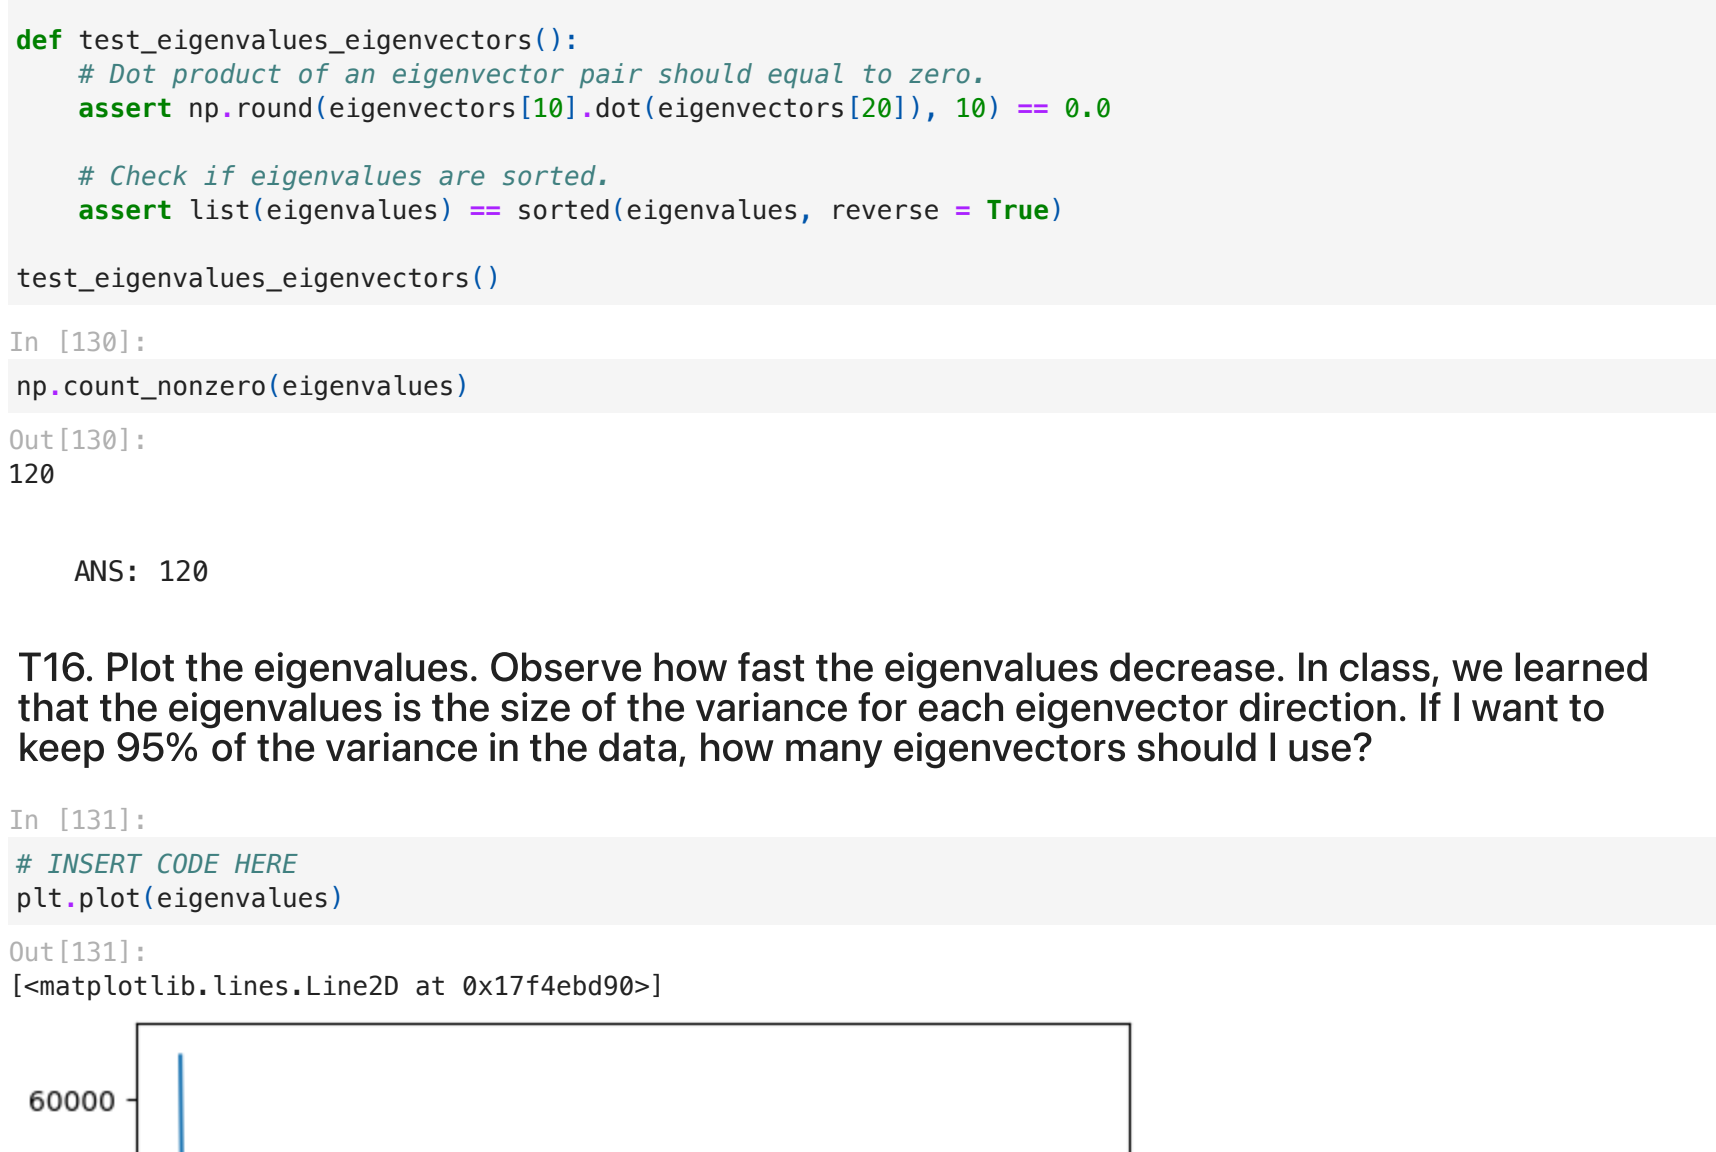
# EER should be either 0.9871428571428571 or 0.9183759388496248 depending on method.
# Recall rate at 0.1% false alarm rate should be 0.5428571428571428.

ANS: EER is where false alarm rate and false negative rate are equal.
```

T11. Compute the mean vector from the training images. Show the vector as an image (use numpy.reshape()). This is typically called the meanface (or meanvoice for speech signals). You answer should look exactly like the image shown below.

```
In [183]:
# INSERT CODE HERE
meanface = np.mean(organize_shape(T), axis=0).reshape(56, 46)
plt.imshow(meanface, cmap="gray")

Out[183]:
<matplotlib.image.AxesImage at 0x177f86858>
```



T12. What is the size of the covariance matrix? What is the rank of the covariance matrix?

```
In [157]:
# T000: Find the size and the rank of the covariance matrix.
X = T.reshape(120, -1).T
X_hat = X - meanface.reshape(-1, 1)

cov_matrix = np.cov(organize_shape(T).T)
print(cov_matrix.shape)
print(np.linalg.matrix_rank(cov_matrix))
(2576, 2576)
119

In [ ]:

ANS: (2576, 2576), 119
```

T13. What is the size of the Gram matrix? What is the rank of Gram matrix? If we compute the eigenvalues from the Gram matrix, how many non- zero eigenvalues do we expect to get?

```
In [187]:
# T000: Compute gram matrix.
gram_matrix = np.matmul(organize_shape(T), organize_shape(T).T)

In [188]:
# T000: Show size and rank of Gram matrix.
print(gram_matrix.shape)
print(np.linalg.matrix_rank(gram_matrix))
(120, 120)
120

ANS: (120, 120), 120
```

T14. Is the Gram matrix also symmetric? Why?

ANS: Yes, because the input is symmetric, so inner product of the input is also symmetric.

T15. Compute the eigenvectors and eigenvalues of the Gram matrix, v and λ . Sort the eigenvalues and eigenvectors in descending order so that the first eigenvalue is the highest, and the first eigenvector corresponds to the best direction. How many non-zero eigenvalues are there? If you see a very small value, it is just numerical error and should be treated as zero.

```
In [129]:
# Hint: https://numpy.org/doc/stable/reference/generated/numpy.linalg.eigh.html

def calculate_eigenvectors_and_eigenvalues(matrix):
    """
    T000: Calculate eigenvectors and eigenvalues,
    then sort the eigenvalues and eigenvectors in descending order.
    Hint: https://numpy.org/doc/stable/reference/generated/numpy.linalg.eigh.html
    """
    # INSERT CODE HERE
    eigenvalues, eigenvectors = np.linalg.eigh(matrix)

    # sort eigenvalues and eigenvectors in descending order
    idx = eigenvalues.argsort()[::-1]
    eigenvalues = eigenvalues[idx]
    eigenvectors = eigenvectors[:, idx]

    return eigenvalues, eigenvectors

eigenvalues, eigenvectors = calculate_eigenvectors_and_eigenvalues(gram_matrix)

def test_eigenvalues_eigenvectors():
    # Dot product of an eigenvector pair should equal to zero.
    assert np.round(eigenvectors[10].dot(eigenvectors[20]), 10) == 0.0

    # Check if eigenvalues are sorted.
    assert list(eigenvalues) == sorted(eigenvalues, reverse = True)

test_eigenvalues_eigenvectors()
```

```
In [130]:
np.count_nonzero(eigenvalues)

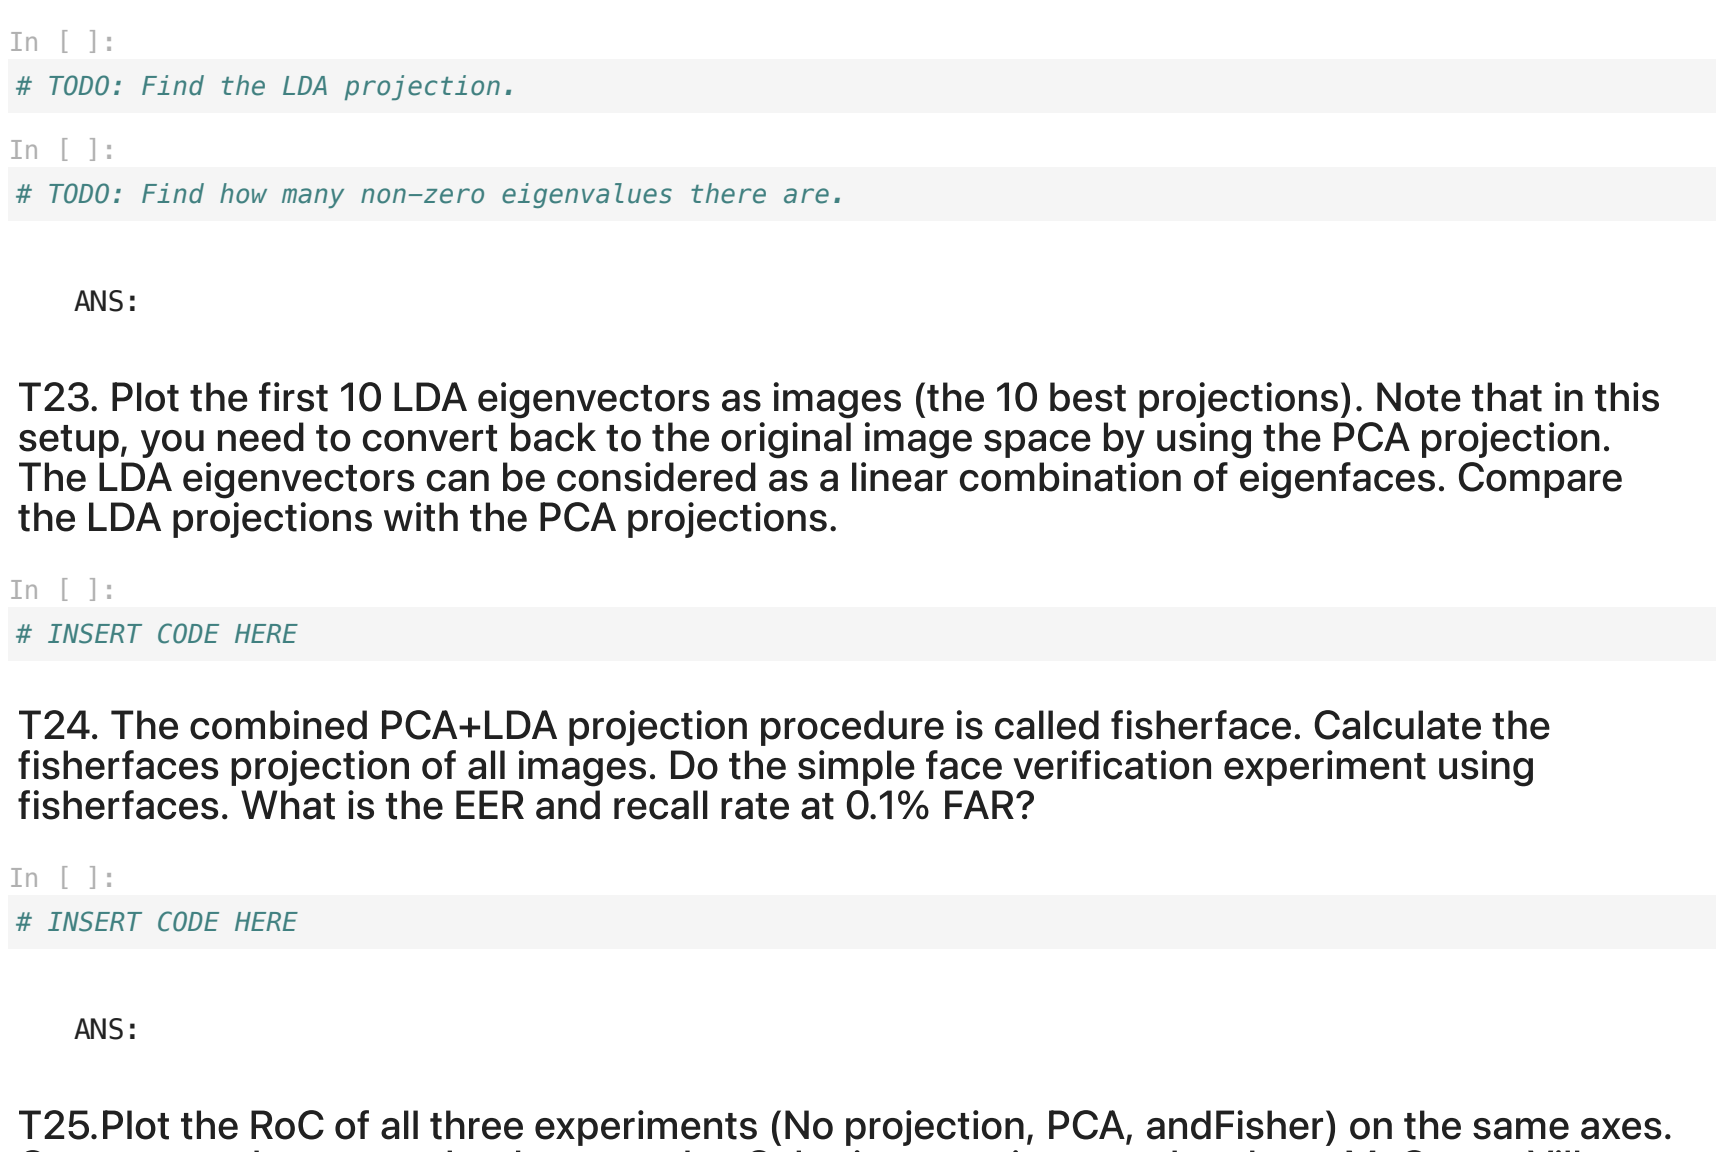
Out[130]:
120

ANS: 120
```

T16. Plot the eigenvalues. Observe how fast the eigenvalues decrease. In class, we learned that the eigenvalues is the size of the variance for each eigenvector direction. If I want to keep 95% of the variance in the data, how many eigenvectors should I use?

```
In [131]:
# INSERT CODE HERE
plt.plot(eigenvalues)

Out[131]:
<matplotlib.lines.Line2D at 0x174ebd98>
```



T17. Compute $\{vec\{v\}$. Don't forget to renormalize so that the norm of each vector is 1 (you can use numpy.linalg.norm). Show the first 10 eigenvectors as images. Two example eigenvectors are shown below. We call these images eigenfaces (or eigenvoice for speech signals).

```
In [153]:
# T000: Compute v, then renormalize it.
# INSERT CODE HERE

v = np.matmul(organize_shape(T).T, eigenvectors[:, 0])
v = v / np.linalg.norm(v)

In [154]:
def test_eigenvector_cov_norm(v):
    assert np.round(np.linalg.norm(v, axis=0), 1) == 1.0).all()

test_eigenvector_cov_norm(v)

In [155]:
# T000: Show the first 10 eigenvectors as images.

v = []
for ev in eigenvectors.T:
    v.append(np.dot(X_hat, ev))

v = np.array(v) / np.linalg.norm(v, axis=1).reshape(-1, 1)

fig, axes = plt.subplots(2, 5, figsize=(10, 5))
for i in range(10):
    eigenface = v[i].reshape(56, 46)
    axes[i // 5, i // 5].imshow(eigenface, cmap="gray")
```



T18. From the image, what do you think the first eigenvector captures? What about the second eigenvector? Look at the original images, do you think biggest variance are capture in these two eigenvectors?

ANS: I think the first eigenvector captures the brightness of the image. The second eigenvector captures the contrast between the face and the hair. I don't think the biggest variance are captured in these two eigenvectors.

T19. Find the projection values of all images. Keep the first $k = 10$ projection values. Repeat the simple face verification system we did earlier using these projected values. What is the EER and the recall rate at 0.1% FAR?

```
In [162]:
def calculate_projection_vectors(matrix, meanface, v):
    """
    T000: Find the projection vectors on v from given matrix and meanface.
    """
    # INSERT CODE HERE
    X = matrix.reshape(120, -1).T
    X_hat = X - meanface.reshape(-1, 1)
    projection_vectors = np.dot(X_hat, v.T)

    return projection_vectors

In [163]:
# T000: Get projection vectors of T and D, then Keep first k projection values.
k = 10
T_reduced = calculate_projection_vectors(T, meanface, v[1:,k])
D_reduced = calculate_projection_vectors(D, meanface, v[1:,k])

def test_reduce_dimension():
    assert T_reduced.shape[-1] == k
    assert D_reduced.shape[-1] == k

test_reduce_dimension()
```

```
ValueError                                Traceback (most recent call last)
Cell In[163], line 4
      2 k = 10
      3 T_reduced = calculate_projection_vectors(T, meanface, v[1:,k])
----> 4 D_reduced = calculate_projection_vectors(D, meanface, v[1:,k])
      7 def test_reduce_dimension():
      8     assert T_reduced.shape[-1] == k

Cell In[162], line 7, in calculate_projection_vectors(matrix, meanface, v)
      2 """
      3 T000: Find the projection vectors on v from given matrix and meanface.
      4 """
      5 # INSERT CODE HERE
----> 6 X = matrix.reshape(120, -1).T
      8 X_hat = X - meanface.reshape(-1, 1)
      9 projection_vectors = np.dot(X_hat, v.T)

ValueError: cannot reshape array of size 721280 into shape (120,newaxis)
```

```
In [ ]:
# T000: Get similarity matrix of T_reduced and D_reduced

In [ ]:
# T000: Find EER and the recall rate at 0.1% FAR.

ANS:
```

T20. What is the k that gives the best EER? Try $k = 5, 6, 7, 8, 9, 10, 11, 12, 13, 14$.

```
In [ ]:
# INSERT CODE HERE

ANS:
```

T21. In order to assure that S_W is invertible we need to make sure that S_W is full rank. How many PCA dimensions do we need to keep in order for S_W to be full rank? (Hint: How many dimensions does S_W have? In order to be of full rank, you need to have the same number of linearly independent factors)

```
In [ ]:
ANS:

# T000: Define dimension of PCA.
n_dim = ...

# T000: Find PCA of T and D with n_dim dimension.
```

T22. Using the answer to the previous question, project the original input to the PCA subspace. Find the LDA projections. To find the inverse, use $-1 \cdot \text{numpy.linalg.pinv}$. Is $S_W S_B$ symmetric? Can we still use numpy.linalg.eigh? How many non-zero eigenvalues are there?

```
In [ ]:
# T000: Find the LDA projection.

In [ ]:
# T000: Find how many non-zero eigenvalues there are.

ANS:
```

T23. Plot the first 10 LDA eigenvectors as images (the 10 best projections). Note that in this setup, you need to convert back to the original image space by using the PCA projection. The LDA eigenvectors can be considered as a linear combination of eigenfaces. Compare the LDA projections with the PCA projections.

```
In [ ]:
# INSERT CODE HERE
```

T24. The combined PCA+LDA projection procedure is called fisherface. Calculate the fisherfaces projection of all images. Do the simple face verification experiment using fisherfaces. What is the EER and recall rate at 0.1% FAR?

```
In [ ]:
# INSERT CODE HERE

ANS:
```

T25. Plot the RoC of all three experiments (No projection, PCA, and Fisher) on the same axes. Compare and contrast the three results. Submit your writeup and code on MyCourseVille.

```
In [ ]:
# INSERT CODE HERE

ANS:
```